



# Graph Analysis Team Introduction

## Experience in GraphBLAS and Graph Databases

Semyon Grigorev

Saint Petersburg State University

April 30, 2022

- Established in 2012
- Lead: Semyon Grigorev
  - ▶ PhD (2016), Associate professor (2016, SPbSU)
  - ▶ dblp: <https://dblp.org/pid/181/9903.html>
  - ▶ h-index (scopus): 5
  - ▶ s.v.grigoriev@spbu.ru
- PhD students: 2
- Master students: 5
- Bachelor students: 6
- Publications
  - ▶ Total: > 30
  - ▶ Scopus: 25

- High-performance graph analysis
  - ▶ GraphBLAS-based algorithms design, implementation and evaluation
  - ▶ Portable multi-GPGPU implementation of GraphBALS-like API
  - ▶ GraphBLAS API analysis

- High-performance graph analysis
  - ▶ GraphBLAS-based algorithms design, implementation and evaluation
  - ▶ Portable multi-GPGPU implementation of GraphBALS-like API
  - ▶ GraphBLAS API analysis
- Formal Language Constrained Path Querying (FLPQ)
  - ▶ New algorithms development
  - ▶ Complexity analysis
  - ▶ New classes of languages investigation
  - ▶ High performance algorithms implementation and evaluation

# GraphBLAS API

- Graph-matrix duality
- Operations over matrices and vectors
  - ▶ Parametrized by semiring-like structures
  - ▶ Based on sparse data structures
  - ▶ Highly parallel

---

<sup>1</sup><https://github.com/DrTimothyAldenDavis/GraphBLAS>

<sup>2</sup><https://github.com/gunrock/graphblast>

<sup>3</sup><https://gitee.com/CSL-ALP/graphblas>

# GraphBLAS API

- Graph-matrix duality
- Operations over matrices and vectors
  - ▶ Parametrized by semiring-like structures
  - ▶ Based on sparse data structures
  - ▶ Highly parallel
- High-performance implementations
  - ▶ SuiteSparse:GraphBLAS<sup>1</sup>: pure C
  - ▶ GraphBLAST<sup>2</sup>: GPGPU, Cuda C
  - ▶ Huawei's GraphBLAS<sup>3</sup>: C++
  - ▶ ...

---

<sup>1</sup><https://github.com/DrTimothyAldenDavis/GraphBLAS>

<sup>2</sup><https://github.com/gunrock/graphblast>

<sup>3</sup><https://gitee.com/CSL-ALP/graphblas>

# GraphBLAS API

- Graph-matrix duality
- Operations over matrices and vectors
  - ▶ Parametrized by semiring-like structures
  - ▶ Based on sparse data structures
  - ▶ Highly parallel
- More information on GraphBLAS
  - ▶ Home page: <https://graphblas.org/>
  - ▶ GraphBLAS-related resources: <https://graphblas.org/GraphBLAS-Pointers/>
  - ▶ Introduction to GraphBLAS:  
<http://mit.bme.hu/~szarnyas/grb/graphblas-introduction.pdf>
- High-performance implementations
  - ▶ SuiteSparse:GraphBLAS<sup>1</sup>: pure C
  - ▶ GraphBLAST<sup>2</sup>: GPGPU, Cuda C
  - ▶ Huawei's GraphBLAS<sup>3</sup>: C++
  - ▶ ...

---

<sup>1</sup><https://github.com/DrTimothyAldenDavis/GraphBLAS>

<sup>2</sup><https://github.com/gunrock/graphblast>

<sup>3</sup><https://gitee.com/CSL-ALP/graphblas>

- BFS-like algorithms
  - ▶ BFS: levels, parents, multiple sources
  - ▶ SSSP
  - ▶ ...
- Graph clustering
- Transitive closure based algorithms
  - ▶ APSP
  - ▶ ...
- Triangle counting
- ...



- BFS-like algorithms
  - ▶ BFS: levels, parents, multiple sources
  - ▶ SSSP
  - ▶ ...
- Graph clustering
- Transitive closure based algorithms
  - ▶ APSP
  - ▶ ...
- Triangle counting
- ...
- LAGraph: collection of GraphBLAS-based algorithms
  - ▶ GitHub: <https://github.com/GraphBLAS/LAGraph>
  - ▶ Latest report: <https://arxiv.org/pdf/2104.01661.pdf>

# Formal Language Constrained Path Querying (FLPQ)

- A way to use formal languages to specify constraints on paths
  - ▶ Regular path querying (RPQ): regular expressions as constraints
  - ▶ Context-free path querying (CFPQ): context-free grammars as constraints

# Formal Language Constrained Path Querying (FLPQ)

- A way to use formal languages to specify constraints on paths
  - ▶ Regular path querying (RPQ): regular expressions as constraints
  - ▶ Context-free path querying (CFPQ): context-free grammars as constraints
- Applications
  - ▶ Graph analysis
  - ▶ Interprocedural static code analysis
  - ▶ Graph database querying

# Our Results: GraphBLAS Implementation

- Tools

- ▶ SPbLA: library of GPGPU-powered sparse boolean linear algebra operations
- ▶ Spla: sparse linear algebra framework for multi-GPU computations based on OpenCL

# Our Results: GraphBLAS Implementation


- Tools

- ▶ SPbLA: library of GPGPU-powered sparse boolean linear algebra operations
- ▶ Spla: sparse linear algebra framework for multi-GPU computations based on OpenCL

- Papers


- ▶ SPbLA: The Library of GPGPU-Powered Sparse Boolean Linear Algebra Operations (GrAPL@IPDPS)

# GraphBLAS Implementation Evaluation: SPbLA<sup>1</sup>



<sup>1</sup>More details: <https://github.com/JetBrains-Research/spbla#performance>

# GraphBLAS Implementation Evaluation: SPbLA<sup>1</sup>



<sup>1</sup>More details: <https://github.com/JetBrains-Research/spbla#performance>

# GraphBLAS Implementation Evaluation: Spla

**Table:** Triangles counting algorithms evaluation results. Time in milliseconds (lower is better).

| Dataset           |        |       | Nvidia |        |        | Intel   |         |
|-------------------|--------|-------|--------|--------|--------|---------|---------|
| Name              | #V     | #E    | GR     | GB     | SP     | SS      | SP      |
| coAuthorsCiteseer | 227.3K | 1.6M  | 2.1    | 2.0    | 9.5    | 17.5    | 64.9    |
| coPapersDBLP      | 540.4K | 30.4M | 5.7    | 94.4   | 201.9  | 543.1   | 1537.8  |
| roadNet-CA        | 1.9M   | 5.5M  | 34.3   | 5.8    | 16.1   | 47.1    | 357.6   |
| com-Orkut         | 3M     | 234M  | 218.1  | 1583.8 | 2407.4 | 23731.4 | 15049.5 |
| cit-Patents       | 3.7M   | 16.5M | 49.7   | 52.9   | 90.6   | 698.3   | 684.1   |
| soc-LiveJournal   | 4.8M   | 68.9M | 69.1   | 449.6  | 673.9  | 4002.6  | 3823.9  |

Tools: Gunrock (GR), GraphBLAST (GB), SuiteSparse (SS), Spla (SP)



# GraphBLAS Implementation Evaluation: Spla

**Table:** Triangles counting algorithms evaluation results. Time in milliseconds (lower is better).

| Dataset           |        |       | Nvidia |        |        | Intel   |         |
|-------------------|--------|-------|--------|--------|--------|---------|---------|
| Name              | #V     | #E    | GR     | GB     | SP     | SS      | SP      |
| coAuthorsCiteseer | 227.3K | 1.6M  | 2.1    | 2.0    | 9.5    | 17.5    | 64.9    |
| coPapersDBLP      | 540.4K | 30.4M | 5.7    | 94.4   | 201.9  | 543.1   | 1537.8  |
| roadNet-CA        | 1.9M   | 5.5M  | 34.3   | 5.8    | 16.1   | 47.1    | 357.6   |
| com-Orkut         | 3M     | 234M  | 218.1  | 1583.8 | 2407.4 | 23731.4 | 15049.5 |
| cit-Patents       | 3.7M   | 16.5M | 49.7   | 52.9   | 90.6   | 698.3   | 684.1   |
| soc-LiveJournal   | 4.8M   | 68.9M | 69.1   | 449.6  | 673.9  | 4002.6  | 3823.9  |

Tools: Gunrock (GR), GraphBLAST (GB), SuiteSparse (SS), Spla (SP)

# GraphBLAS Implementation Evaluation: Spla

**Table:** Triangles counting algorithms evaluation results. Time in milliseconds (lower is better).

| Dataset           |        |       | Nvidia |        |        | Intel   |         |
|-------------------|--------|-------|--------|--------|--------|---------|---------|
| Name              | #V     | #E    | GR     | GB     | SP     | SS      | SP      |
| coAuthorsCiteseer | 227.3K | 1.6M  | 2.1    | 2.0    | 9.5    | 17.5    | 64.9    |
| coPapersDBLP      | 540.4K | 30.4M | 5.7    | 94.4   | 201.9  | 543.1   | 1537.8  |
| roadNet-CA        | 1.9M   | 5.5M  | 34.3   | 5.8    | 16.1   | 47.1    | 357.6   |
| com-Orkut         | 3M     | 234M  | 218.1  | 1583.8 | 2407.4 | 23731.4 | 15049.5 |
| cit-Patents       | 3.7M   | 16.5M | 49.7   | 52.9   | 90.6   | 698.3   | 684.1   |
| soc-LiveJournal   | 4.8M   | 68.9M | 69.1   | 449.6  | 673.9  | 4002.6  | 3823.9  |

Tools: Gunrock (GR), GraphBLAST (GB), SuiteSparse (SS), Spla (SP)

# GraphBLAS Implementation Evaluation: Spla

**Table:** Triangles counting algorithms evaluation results. Time in milliseconds (lower is better).

| Dataset           |        |       | Nvidia |        |        | Intel   |         |
|-------------------|--------|-------|--------|--------|--------|---------|---------|
| Name              | #V     | #E    | GR     | GB     | SP     | SS      | SP      |
| coAuthorsCiteseer | 227.3K | 1.6M  | 2.1    | 2.0    | 9.5    | 17.5    | 64.9    |
| coPapersDBLP      | 540.4K | 30.4M | 5.7    | 94.4   | 201.9  | 543.1   | 1537.8  |
| roadNet-CA        | 1.9M   | 5.5M  | 34.3   | 5.8    | 16.1   | 47.1    | 357.6   |
| com-Orkut         | 3M     | 234M  | 218.1  | 1583.8 | 2407.4 | 23731.4 | 15049.5 |
| cit-Patents       | 3.7M   | 16.5M | 49.7   | 52.9   | 90.6   | 698.3   | 684.1   |
| soc-LiveJournal   | 4.8M   | 68.9M | 69.1   | 449.6  | 673.9  | 4002.6  | 3823.9  |

Tools: Gunrock (GR), GraphBLAST (GB), SuiteSparse (SS), Spla (SP)

# Our Results: Graph Databases Extensions with CFPQ

- Tools

- ▶ GLL4Graph: CFPQ for Neo4j
- ▶ CFPQ for RedisGraph

# Our Results: Graph Databases Extensions with CFPQ


- Tools

- ▶ GLL4Graph: CFPQ for Neo4j
- ▶ CFPQ for RedisGraph

- Papers

- ▶ Multiple-Source Context-Free Path Querying in Terms of Linear Algebra (EDBT, Core A)
- ▶ Context-free path querying by matrix multiplication (GRADES-NDA@SIGMOD)

# Graph Databases Extensions with CFPQ: Evaluation



# Our Results: FLPQ Algorithms Evaluation

- Tools

- ▶ CFPQ\_PyAlgo: set of GraphBLAS-based FLPQ algorithms
  - ✓ GraphBLAS-based CFPQ algorithms
  - ⚙ GraphBLAS-based RPQ algorithms
  - ⚙ GraphBLAS-based MCFG-PQ algorithms

# Our Results: FLPQ Algorithms Evaluation

- Tools

- ▶ CFPQ\_PyAlgo: set of GraphBLAS-based FLPQ algorithms
  - ✓ GraphBLAS-based CFPQ algorithms
  - ⚙ GraphBLAS-based RPQ algorithms
  - ⚙ GraphBLAS-based MCFG-PQ algorithms
- ▶ CFPQ\_Data: dataset for FLPQ algorithms evaluation
  - ✓ Graphs for CFPQ evaluation
  - ⚙ Graphs for RPQ evaluation
  - ⚙ Integration with LDBC Graphalytics



# Our Results: FLPQ Algorithms Evaluation

- Tools

- ▶ CFPQ\_PyAlgo: set of GraphBLAS-based FLPQ algorithms
  - ✓ GraphBLAS-based CFPQ algorithms
  - ⚙ GraphBLAS-based RPQ algorithms
  - ⚙ GraphBLAS-based MCFG-PQ algorithms
- ▶ CFPQ\_Data: dataset for FLPQ algorithms evaluation
  - ✓ Graphs for CFPQ evaluation
  - ⚙ Graphs for RPQ evaluation
  - ⚙ Integration with LDBC Graphalytics
- ⚙ LDBC Graphalytics extension for evaluation of formal language constrained path querying

# Our Results: FLPQ Algorithms Evaluation

- Tools

- ▶ CFPQ\_PyAlgo: set of GraphBLAS-based FLPQ algorithms

- ✓ GraphBLAS-based CFPQ algorithms

- ⚙ GraphBLAS-based RPQ algorithms

- ⚙ GraphBLAS-based MCFG-PQ algorithms

- ▶ CFPQ\_Data: dataset for FLPQ algorithms evaluation

- ✓ Graphs for CFPQ evaluation

- ⚙ Graphs for RPQ evaluation

- ⚙ Integration with LDBC Graphalytics

- ⚙ LDBC Graphalytics extension for evaluation of formal language constrained path querying

- Papers

- ▶ Evaluation of the context-free path querying algorithm based on matrix multiplication (GRADES-NDA@SIGMOD)