SPLA: Portable Generic Sparse Linear Algebra Library for GPU Computations

1st Egor Orachev
Faculty of Mathematics and Mechanics
St. Petersburg State University,
Programming Languages and Tools Lab
JetBrains Research
St. Petersburg, Russia
egor.orachev@gmail.com
0000-0002-0424-4059

2nd Gleb Marin
School of Physics, Mathematics,
and Computer Science
HSE University
St. Petersburg, Russia
glebmar2001@gmail.com
0000-0002-0873-1647

3rd Semyon Grigorev
Faculty of Mathematics and Mechanics
St. Petersburg State University,
Programming Languages and Tools Lab
JetBrains Research
St. Petersburg, Russia
s.v.grigoriev@spbu.ru
semyon.grigorev@jetbrains.com
0000-0002-7966-0698

Abstract—Scalable high-performance graph analysis is an actual nontrivial challenge and usage of sparse linear algebra operations as building blocks for graph analysis algorithms, which is a core idea of GraphBLAS standard, is a promising way to attack it. While it is known that sparse linear algebra operations can be efficiently implemented on GPGPU, full GraphBLAS implementation on GPGPU is a nontrivial task that is almost solved by GraphBLAST project. Though it is shown by GraphBLAST that utilization of GPGPUs for GraphBLAS implementation significantly improves performance, portability and scalability problems are not solved yet: GraphBLAST uses Nvidia stack and utilizes only one GPGPU. In this work we propose an SPLA library that aimed to solve these problems: it uses OpenCL to be portable and designed to utilize multiple GPGPUs. Preliminary evaluation shows that !!!!

Index Terms—graphs, algorithms, graph analysis, sparse linear algebra, GraphBLAS, GPGPU, OpenCL

I. INTRODUCTION

Scalable high-performance graph analysis is an actual challenge. There is a big number of ways to attack this challenge [1] and the first promising idea is to utilize general-purpose graphic processing units (GPGPU-s). Such existing solutions, as CuSha [2] and Gunrock [3] show that utilization of GPUs can improve the performance of graph analysis, moreover it is shown that solutions may be scaled to multi-GPU systems. But low flexibility and high complexity of API are problems of these solutions.

The second promising thing which provides a user-friendly API for high-performance graph analysis algorithms creation is a GraphBLAS API [4] which provides linear algebra based building blocks to create graph analysis algorithms. The idea of GraphBLAS is based on is a well-known fact that linear algebra operations can be efficiently implemented on parallel hardware. Along with this, a graph can be natively represented using matrices: adjacency matrix, incidence matrix, etc. While reference CPU-based implementation of GraphBLAS, SuiteSparse:GraphBLAS [5], demonstrates good performance in real-world tasks, GPU-based implementation is challenging.

One of the challenges in this way is that real data are often sparse, thus underlying matrices and vectors are also sparse, and, as a result, classical dense data structures and respective algorithms are inefficient. So, it is necessary to use advanced data structures and procedures to implement sparse linear algebra, but the efficient implementation of them on GPU is hard due to the irregularity of workload and data access patterns. Though such well-known libraries as cuSparse show that sparse linear algebra operations can be efficiently implemented for GPGPU-s, it is not so trivial to implement GraphBLAS on GPGPU. First of all, it requires *generic* sparse linear algebra, thus it is impossible just to reuse existing libraries which are almost all specified for operations over floats. The second problem is specific optimizations, such as maskings fusion, which can not be natively implemented on top of existing kernels. Nevertheless, there is a number of implementations of GraphBLAS on GPGPU, such as Graph-BLAST: [6], GBTL [7], which show that GPGPUs utilization can improve the performance of GraphBLAS-based graph analysis solutions. But these solutions are not portable because they are based on Nvidia Cuda stack. Moreover, the scalability problem is not solved: all these solutions support only single-GPU, not multi-GPU computations.

To provide portable multi-GPU implementation of Graph-BLAS API we developed a *SPLA* library (sources are published on GitHub: https://github.com/JetBrains-Research/spla). This library utilizes OpenCL for GPGPU computing to be portable across devices of different vendors. Moreover, it is initially designed to utilize multiple GPGPUs to be scalable. To sum up, the contribution of this work is the following.

- Design of portable multi-GPU GraphBLAS implementation proposed. Additionally, proposed design is aimed to simplify library tuning and wrappers for different high-level platforms and languages creation.
- Subset of GraphBLAS API, including such operations as masking, matrix-matrix multiplication, matrix-matrix e-wise addition, is implemented. Current implementation is limited by COO and CSR matrix representation format and uses basic algorithms for some operations, but work in progress and more data formats will be supported and

advanced algorithms will be implemented in the future.

- Preliminary evaluation on such algorithms as breadth-first search (BFS) and triangles counting (TC), and real-world graphs shows portability across different vendors and promising performance: for some problems Spla is comparable with GraphBLAST. Surprisingly, for some problems, the proposed solution on embedded Intel graphic card shows better performance than SuiteSparse:GraphBLAS on the same CPU. At the same time, the evaluation shows that further optimization is required.

II. SOLUTION DESCRIPTION

A. Design Principles

SPLA library is designed in the way to maximize potential library performance, simplify its implementation and extensions, and to provided the end-user verbose, but effective interface allowing customization and precise control over operations execution. These ideas are captured in the following principles.

- *DAG-based expressions*. User constructs a computational expression from basic nodes and uses oriented edges to describe data dependencies between these nodes.
- *Automated hybrid-storage format*. Library uses internally specialized preprocessing to format data and automate its sharing between computational nodes.
- *Automated multi-GPU scheduling*. Computational work is automatically scheduled between available devices for execution. Scheduling order, dependencies and granularity are defined from DAG expression, submitted by user.
- *Customization of primitive types and operations*. Underlying primitives types and functions can be customized by user. Customization process does not requires library re-compilation.
- *Exportable interface*. Library has C++ interface with an automated reference-counting and with no-templates usage. It can be wrapped by C99 compatible API and exported to other languages, for example, in a form of a Python-package.

B. Architecture Overview

Library general execution architecture is depicted in Fig. 1. As an input library accepts expression composed in the form of a DAG. Nodes represent fundamental operations, such as matrix-matrix multiplication. Links describe dependencies between nodes. Expression execution is *asynchronous*. User can block and wait until its completion, or without blocking probe the expression until it is either *completed* or *aborted*.

Expression is transformed into a task graph. Task graph is submitted for execution to the task manager. Each task is processed by specialised *NodeProcessor*, capable of processing particular node type. Each task, when executed, is split dynamically into a set of parallel sub-tasks. Each sub-task is processed by specialized *Algorithm*, which is capable of processing input blocks of matrices or vectors in particular storage formats with specific set of options. *NodePorcessor* and *Algorithm* are selected at runtime from a registry using


Fig. 1. Library expression processing architecture.

properties and arguments of the expression. Thus, it allows precise processing and optimization of edge-cases.

Granularity level of sub-tasks is defined by the structure of underlying processed primitives. Target device for execution is automatically assigned for the sub-task based on expression and node parameters. Currently, the fixed and uniform distribution for assignment is supported.

C. Containers

Library provides general *M-by-N Matrix*, *N Vector* and *Scalar* primitives. Underlying primitives types is specified by *Type* object. Internally primitives are stored in a hybrid storage in a form of two- or one- dimensional blocks' grid form matrices and vectors respectively. Each block is empty (not stored) or store some data in any format. Blocks are immutable, they can be safely shared across computational devices.

Currently COO/CSR blocks for matrices, and COO/Dense blocks for vectors a are supported. Format choice is motivated by its simplicity and easy of implementation. Other formats, such as CSC, DCSR, ELL, etc. can be added to the library by implementation of formats convertation or by the specialization of *Algorithm* for a specific format.

D. Algebraic Operations

Library supports all commonly used linear algebra operations, such as *mxm*, *vxm*, *eadd*, *reduce*, *transpose*. Other operations are coming soon, since library still in development. Interface of operations is designed *similar* to GraphBLAS. It supports *masking*, *accum* of the result, *add* and/or *mult* user-functions specification, and *descriptor* object for additional operation tweaking.

E. Implementation Details

Library uses OpenCL 1.2 API as underlying compute API. Boost Compute [8] is utilized as a high-level library on top of the OpenCL functionality. It provides thread-safe kernel caching, meta-kernel programming, and a set of basic parallel primitives such as *device vector*, *sort*, *reduce*, *scan*, etc. which was extended further to meet this project requirements. Taskflow [9] is used as tasking library. It supports task-dependencies and dynamic tasking, utilized in order to create and execute sub-tasks.

User-defined *Types* are represented as POD-structures, and handled by the library as a fixed-size sequences of bytes. User-defined *Functions* are effectively textual strings with OpenCL

code, injected into generalized meta-kernels. Library has a number of predefined types, such as *signed/unsigned integers*, *floating point* types, and a set of common operations, such as *arithmetic*, *logic*, *first/second*, etc.

For particular SpVSpM implementation ESC algorithm [10] is employed. Masked SpGEMM is based on Yang et al. work [6] solution. Tiled GPU merge path [11] algorithm utilized for element-wise addition and masking implementation. The code is generalized and written in a form of meta-kernels, so actual functions for elements reduction or multiplication are injected later. Kernel compilation is done on demand if no previously cached entry present.

III. EVALUATION

For performance analysis of proposed solution we evaluated some most common graph algorithms using real-world sparse matrix data. As a baseline for comparison we chose LAGraph [12] in connection with SuiteSparse [5] as a CPU tool, Gunrock [3] and GraphBLAST [6] as a Nvidia GPU tools. Also, we tested algorithms on several devices with distinct OpenCL vendors in order to validate portability of the proposed solution. In general, these evaluation intentions are summarized in the following research questions.

RQ1 What is the performance of the proposed solution relative to existing tools for both CPU and GPU analysis?

RQ2 What is the portability of the proposed solution with respect to various device vendors and OpenCL runtimes?

A. Evaluation Setup

For evaluation, we use a PC with Ubuntu 20.04 installed, which has 3.40Hz Intel Core i7-6700 4-core CPU, DDR4 64Gb RAM, and Nvidia GeForce GTX 1070 GPU with 8Gb VRAM. Host programs were compiled with GCC 9.3.0 compiler. Programs using CUDA were compiled with GCC 8.4.0 and Nvidia NVCC 10.1.243 compiler. Release mode and maximum optimization level was enabled for all tested programs. Data loading time, preparation, format transformations and host-device initial communications are excluded from time measurements. All tests are averaged across 10 runs. Additional warm-up run for each test execution is excluded from measurements.

B. Graph Algorithms

For preliminary study *breadth-first search* (bfs) and *triangles counting* (tc) algorithms were chosen, since they allows analyse the performance of *vxm* and *mxm* operations, rely heavily on *masking*, and utilize *reduction* or *assignment*. BFS implementation utilizes automated vector storage from sparse to dense switch and only push optimization. TC implementation uses masked *mxm* of source lower-triangular matrix with second transposed argument.

C. Dataset

Nine graph matrices were selected from the Sparse Matrix Collection at University of Florida [13]. Information about graphs is summarized in Table I. All datasets are converted to undirected graphs. Self-loops and duplicated edges are removed.

TABLE I
DATASET DESCRIPTION.

Dataset	Vertices	Edges	Max Degree
coAuthorsCiteseer	227.3K	1.6M	1372
coPapersDBLP	540.4K	30.4M	3299
hollywood-2009	1.1M	113.8M	11,467
roadNet-CA	1.9M	5.5M	12
com-Orkut	3M	234M	33313
cit-Patents	3.7M	16.5M	793
rgg_n_2_22_s0	4.1M	60.7M	36
soc-LiveJournal	4.8M	68.9M	20,333
indochina-2004	7.5M	194.1M	256,425

D. Results

Table II presents results of the evaluation and compares performance of Spla against other tool on different execution platforms. Tools are grouped by the type of the device for the execution, where either Nvidia GPU or Intel CPU are used. Cell left empty if tested tool failed to analyse graph due to *out of memory* exception.

In general, Spla BFS shows acceptable performance, especially on graphs with large vertex degree, such as soc-LiveJournal and com-Orkut. On graphs roadNet-CA and rgg it has a significant performance drop due to the nature of underlying algorithms and data structures. Firstly, library utilizes immutable data buffers. Thus, iteratively updated dense vector of reached vertices must be copied for each modification, what dominates the performance of the library on a graph with large search depth. Secondly, Spla BFS does not utilise *pull optimization*, what is critical in a graph with relatively small search frontier.

Spla TC has a good performance on GPU, which is better in all cases that reference SuiteSparse solution. But in most tests GPU competitors, especially Gunrock, show smaller processing times. GraphBLAST shows better performance as well. Library utilises masked SpGEMM algorithm, the same as in GraphBLAST, but without *identity* element to fill gaps. Library explicitly stores all non-zero elements, and uses mask to reduce only non-zero while evaluating dot products of rows and columns. What causes extra divergence inside work groups. On Intel device Spla shows better performance compared to SuiteSparse on com-Orkut, cit-Patents and soc-LiveJournal. A possible reason is the large lengths of processed rows and columns in the product of matrices.

Gunrock shows nearly best average performance due to its specialized and optimized algorithms. Also, it has good time characteristics on a mentioned earlier roadNet-CA and rgg in BFS algortihm. GraphBLAST follows Gunrock and show good performance as well. But it runs out of memory on a two significantly large graphs con-Orkut and indochina-2004. Spla does not rut out of memory on any test due to simplified storage scheme.

TABLE II

GRAPH ALGORITHMS EVALUATION RESULTS.
TIME IN MILLISECONDS (LOWER IS BETTER).

Dataset	Nvidia			Intel	
	GR	GB	SP	SS	SP
BFS					
hollywood-2009	20.3	82.3	36.9	23.7	303.4
roadNet-CA	33.4	130.8	1456.4	168.2	965.6
soc-LiveJournal	60.9	80.6	90.6	75.2	1206.3
rgg_n_2_22_s0	98.7	414.9	4504.3	1215.7	15630.1
com-Orkut	205.2	– –	117.9	43.2	903.6
indochina-2004	32.7	– –	199.6	227.1	2704.6
TC					
coAuthorsCiteseer	2.1	2.0	9.5	17.5	64.9
coPapersDBLP	5.7	94.4	201.9	543.1	1537.8
roadNet-CA	34.3	5.8	16.1	47.1	357.6
com-Orkut	218.1	1583.8	2407.4	23731.4	15049.5
cit-Patents	49.7	52.9	90.6	698.3	684.1
soc-LiveJournal	69.1	449.6	673.9	4002.6	3823.9

Tools: Gunrock (GR), GraphBLAST (GB), SuiteSparse (SS), Spla (SP).

IV. CONCLUSION

In this paper we present portable generic sparse linear algebra library for GPU computations, which provides sparse primitives and supports set of common operations. Evaluation of proposed solution for real-world graph analysis shows, that initial OpenCL-based operations implementation with limited set of optimizations has promising performance compared to other tools and can be easily on devices of multiple vendors, what gives significant flexibility in choose of HPC hardware. Since project still in the active development, the following major and important development tasks must be highlighted.

- *Operations.* New mathematical operations must be implemented, as well as required and most used data manipulation, row/column management, filtering, selection and transformation subroutines must be supported.
- *Performance tuning.* State of the art high performance algorithms must be implemented for such operations, as SpMSpV, SpMV, SpGEMM, etc. in order to increase library efficiency. Also, different techniques, such as automated storage format transitions, pull-push direction optimization, etc. must be employed for better performance on a large range of data.
- *Graph algorithms.* Algorithms such as *single-source shortest paths*, *page rank*, *connected components*, etc. must be implemented using library API and their performance must be analysed as well.
- *API exporting.* Library interface must be exported to other programming languages, such as C or Python. Finally, Python-package for applied graph analysis must be published.
- *Multi-GPU scheduling.* Currently library supports scheduling only on single GPU. But, its internal hybrid storage format and tasking allow multi-node execution. Thus, multi-GPU scheduling and precise data sharing must be implemented and performance analysis of the solution must be carried on.

Finally, we plan to study the characteristics of the implemented library with all algorithms and improvements, and compare its performance to existing tools using uniform benchmarking platform such as a Graphalytics [14].

REFERENCES

[1] M. E. Coimbra, A. P. Francisco, and L. Veiga, "An analysis of the graph processing landscape," *Journal of Big Data*, vol. 8, no. 1, Apr. 2021. [Online]. Available: https://doi.org/10.1186/s40537-021-00443-9

[2] F. Khorasani, K. Vora, R. Gupta, and L. N. Bhuyan, "Cusha: Vertex-centric graph processing on gpus," in *Proceedings of the 23rd International Symposium on High-Performance Parallel and Distributed Computing*, ser. HPDC '14. New York, NY, USA: Association for Computing Machinery, 2014, p. 239–252. [Online]. Available: https://doi.org/10.1145/2600212.2600227

[3] Y. Pan, Y. Wang, Y. Wu, C. Yang, and J. D. Owens, "Multi-gpu graph analytics," in *2017 IEEE International Parallel and Distributed Processing Symposium (IPDPS)*, 2017, pp. 479–490.

[4] J. Kepner, P. Aaltonen, D. Bader, A. Buluç, F. Franchetti, J. Gilbert, D. Hutchison, M. Kumar, A. Lumsdaine, H. Meyerhenke, S. McMillan, C. Yang, J. D. Owens, M. Zalewski, T. Mattson, and J. Moreira, "Mathematical foundations of the graphblas," in *2016 IEEE High Performance Extreme Computing Conference (HPEC)*, 2016, pp. 1–9.

[5] T. A. Davis, "Algorithm 1000: Suitesparse:graphblas: Graph algorithms in the language of sparse linear algebra," *ACM Trans. Math. Softw.*, vol. 45, no. 4, Dec. 2019. [Online]. Available: https://doi.org/10.1145/3322125

[6] C. Yang, A. Buluc, and J. D. Owens, "Graphblast: A high-performance linear algebra-based graph framework on the gpu," 2019.

[7] P. Zhang, M. Zalewski, A. Lumsdaine, S. Misurda, and S. McMillan, "Gbtl-cuda: Graph algorithms and primitives for gpus," in *2016 IEEE International Parallel and Distributed Processing Symposium Workshops (IPDPSW)*, 2016, pp. 912–920.

[8] J. Szuppe, "Boost.compute: A parallel computing library for c++ based on opencl," in *Proceedings of the 4th International Workshop on OpenCL*, ser. IWOCL '16. New York, NY, USA: Association for Computing Machinery, 2016. [Online]. Available: https://doi.org/10.1145/2909437.2909454

[9] T.-W. Huang, D.-L. Lin, C.-X. Lin, and Y. Lin, "Taskflow: A lightweight parallel and heterogeneous task graph computing system," *IEEE Transactions on Parallel and Distributed Systems*, vol. 33, pp. 1303–1320, 2022.

[10] S. Dalton, L. Olson, and N. Bell, "Optimizing sparse matrix—matrix multiplication for the gpu," *ACM Trans. Math. Softw.*, vol. 41, no. 4, oct 2015. [Online]. Available: https://doi.org/10.1145/2699470

[11] O. Green, R. McColl, and D. A. Bader, "Gpu merge path: A gpu merging algorithm," in *Proceedings of the 26th ACM International Conference on Supercomputing*, ser. ICS '12. New York, NY, USA: Association for Computing Machinery, 2012, p. 331–340. [Online]. Available: https://doi.org/10.1145/2304576.2304621

[12] G. Szárnyas, D. A. Bader, T. A. Davis, J. Kitchen, T. G. Mattson, S. McMillan, and E. Welch, "Lagraph: Linear algebra, network analysis libraries, and the study of graph algorithms," 2021.

[13] T. A. Davis and Y. Hu, "The university of florida sparse matrix collection," *ACM Trans. Math. Softw.*, vol. 38, no. 1, dec 2011. [Online]. Available: https://doi.org/10.1145/2049662.2049663

[14] A. Iosup, A. Musaafir, A. Uta, A. P. Pérez, G. Szárnyas, H. Chafi, I. G. Tănase, L. Nai, M. Anderson, M. Capotă, N. Sundaram, P. Boncz, S. Depner, S. Heldens, T. Manhardt, T. Hegeman, W. L. Ngai, and Y. Xia, "The ldbc graphalytics benchmark," 2021.